

DISPLAYING DATA FROM MULTIPLE TABLES

Cartesian Products

- A Cartesian product is formed when:
 - A join condition is omitted
 - A join condition is invalid
 - All rows in the first table are joined to all rows in the second table
- To avoid a Cartesian product, always include a valid join condition in a WHERE clause.

A Cartesian product is generated if a join condition is omitted. The example in the slide displays employee last name and department name from the EMPLOYEES and DEPARTMENTS tables. Because no WHERE clause has been specified, all rows (20 rows) from the EMPLOYEES table are joined with all rows (8 rows) in the DEPARTMENTS table, thereby generating 160 rows in the output.

```
SELECT last_name, department_name dept_name
FROM employees, departments;
```

| LAST_NAME | DEPT_NAME |
|-----------|----------------|
| King | Administration |
| Kochhar | Administration |
| De Haan | Administration |
| ... | ... |
| Gietz | Contracting |

160 rows selected.

Types of Joins

Oracle Proprietary Joins (8i and prior):

- Equijoin
- Nonequijoin
- Outer join
- Self join

SQL: 1999 Compliant Joins:

- Cross joins
- Natural joins
- Using clause
- Full or two sided outer joins
- Arbitrary join conditions for outer joins

Joining Tables Using Oracle Syntax

Use a join to query data from more than one table.

```
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column1 = table2.column2;
```

- Write the join condition in the WHERE clause.
- Prefix the column name with the table name when the same column name appears in more than one table.

Guidelines

- When writing a `SELECT` statement that joins tables, precede the column name with the table name for clarity and to enhance database access.
- If the same column name appears in more than one table, the column name must be prefixed with the table name.
- To join n tables together, you need a minimum of $n-1$ join conditions. For example, to join four tables, a minimum of three joins is required. This rule may not apply if your table has a concatenated primary key, in which case more than one column is required to uniquely identify each row.

For more information, see *Oracle9i SQL Reference*, “`SELECT`.”

What Is an Equijoin?

EMPLOYEES

| EMPLOYEE_ID | DEPARTMENT_ID |
|-------------|---------------|
| 200 | 10 |
| 201 | 20 |
| 202 | 20 |
| 124 | 50 |
| 141 | 50 |
| 142 | 50 |
| 143 | 50 |
| 144 | 50 |
| 103 | 60 |
| 104 | 60 |
| 107 | 60 |
| 109 | 110 |
| 206 | 110 |

Additional Search Conditions

Using the AND Operator

In addition to the join, you may have criteria for your WHERE clause to restrict the rows under consideration for one or more tables in the join. For example, to display employee Matos' department number and department name, you need an additional condition in the WHERE clause.

```
SELECT last_name, employees.department_id,
       department_name
FROM   employees, departments
WHERE  employees.department_id = departments.department_id
AND    last_name = 'Matos';
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Matos | 50 | Shipping |

Qualifying Ambiguous

Column Names

- Use table prefixes to qualify column names that are in multiple tables.
- Improve performance by using table prefixes.
- Distinguish columns that have identical names but reside in different tables by using column aliases.

Using Table Aliases

- Simplify queries by using table aliases
- Improve performance by using table prefixes

```
SELECT e.employee_id, e.last_name, e.department_id,
       d.department_id, d.location_id
FROM   employees e, departments d
WHERE  e.department_id = d.department_id;
```

Joining More than Two Tables

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID |
|-----------|---------------|
| King | 90 |
| Neer | 90 |
| De Haan | 90 |
| Hunold | 60 |
| Baer | 60 |
| Lorentz | 60 |
| ... | ... |
| Trenn | ... |
| Whalen | 10 |
| Patton | 20 |
| Fey | 20 |
| Higgins | 110 |
| Gietz | 110 |

DEPARTMENTS

| DEPARTMENT_ID | LOCATION_ID |
|---------------|-------------|
| 10 | 1500 |
| 20 | 1800 |
| 30 | 1500 |
| 40 | 1400 |
| 50 | 2500 |
| 60 | 1700 |
| 70 | 1700 |
| 80 | 1700 |

LOCATIONS

| LOCATION_ID | CITY |
|-------------|---------------------|
| 1400 | Scottsdale |
| 1500 | South San Francisco |
| 1700 | Seattle |
| 1800 | Toronto |
| 2500 | Orlando |

8 rows selected

20 rows selected

To join n tables together, you need a minimum of $n-1$ join conditions. For example, to join three tables, a minimum of two joins is required.

```

SELECT e.last_name, d.department_name, l.city
FROM   employees e, departments d, locations l
WHERE  e.department_id = d.department_id
AND    d.location_id = l.location_id;

```

| LAST_NAME | DEPARTMENT_NAME | CITY |
|-----------|-----------------|---------------------|
| Hunold | IT | Southlake |
| Ernst | IT | Southlake |
| Lorentz | IT | Southlake |
| Mourgos | Shipping | South San Francisco |
| Rais | Shipping | South San Francisco |

Nonequijoins

| EMPLOYEES | |
|-----------|--------|
| LAST_NAME | SALARY |
| King | 24000 |
| Kochhar | 17000 |
| De Haan | 17000 |
| Hunold | 9000 |
| Ernst | 8000 |
| Lorentz | 4200 |
| Mourgos | 5800 |
| Rais | 3500 |
| Davies | 3100 |
| Matos | 2600 |
| Vargas | 2500 |
| ... | |
| Higgins | 12000 |
| Gietz | 8300 |

20 rows selected.

| JOB_GRADES | | |
|------------|------------|-------------|
| GRA | LOWEST_SAL | HIGHEST_SAL |
| A | 1000 | 2999 |
| B | 3000 | 5999 |
| C | 6000 | 9999 |
| D | 10000 | 14999 |
| E | 15000 | 24999 |
| F | 25000 | 40000 |

6 rows selected.

← Salary in the EMPLOYEES table must be between lowest salary and highest salary in the JOB_GRADES table.

A nonequijoin is a join condition containing something other than an equality operator.

The relationship between the EMPLOYEES table and the JOB_GRADES table has an example of a nonequijoin. A relationship between the two tables is that the SALARY column in the EMPLOYEES table must be between the values in the LOWEST_SALARY and HIGHEST_SALARY columns of the JOB_GRADES table. The relationship is obtained using an operator other than equals (=).

```

SELECT e.last_name, e.salary, j.grade_level
FROM   employees e, job_grades j
WHERE  e.salary BETWEEN j.lowest_sal AND j.highest_sal;

```

| LAST_NAME | SALARY | GRA |
|-----------|--------|-----|
| Matos | 2600 | A |
| Vargas | 2500 | A |
| Lorentz | 4200 | B |
| Mourgos | 5800 | B |
| Rais | 3500 | B |
| Davies | 3100 | B |
| ... | | |
| Kochhar | 17000 | E |
| De Haan | 17000 | E |

The example in the slide creates a nonequijoin to evaluate an employee's salary grade. The salary must be *between* any pair of the low and high salary ranges.

It is important to note that all employees appear exactly once when this query is executed. No employee is repeated in the list. There are two reasons for this:

- None of the rows in the job grade table contain grades that overlap. That is, the salary value for an employee can lie only between the low salary and high salary values of one of the rows in the salary grade table.
- All of the employees' salaries lie within the limits provided by the job grade table. That is, no employee earns less than the lowest value contained in the `LOWEST_SAL` column or more than the highest value contained in the `HIGHEST_SAL` column.

Note: Other conditions, such as `<=` and `>=` can be used, but `BETWEEN` is the simplest. Remember to specify the low value first and the high value last when using `BETWEEN`.

Table aliases have been specified in the slide example for performance reasons, not because of possible ambiguity.

Outer Joins

DEPARTMENTS

| DEPARTMENT_NAME | DEPARTMENT_ID |
|-----------------|---------------|
| Administration | 10 |
| Marketing | 20 |
| Shipping | 50 |
| IT | 60 |
| Sales | 60 |
| Executive | 90 |
| Accounting | 110 |
| Contracting | 190 |

8 rows selected.

EMPLOYEES

| DEPARTMENT_ID | LAST_NAME |
|---------------|-----------|
| 90 | King |
| 90 | Kochhar |
| 90 | De Haan |
| 60 | Hunold |
| 60 | Ernst |
| 10 | Wu |
| 20 | Hartstein |
| 20 | Fay |
| 110 | Higgins |
| 110 | Gietz |

20 rows selected.



There are no employees in department 190.

Returning Records with No Direct Match with Outer Joins

If a row does not satisfy a join condition, the row will not appear in the query result. For example, in the equijoin condition of EMPLOYEES and DEPARTMENTS tables, employee Grant does not appear because there is no department ID recorded for her in the EMPLOYEES table. Instead of seeing 20 employees in the result set, you see 19 records.

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE e.department_id = d.department_id;
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Whalen | 10 | Administration |
| Hartstein | 20 | Marketing |
| Fay | 20 | Marketing |
| Mourgos | 50 | Shipping |
| ... | ... | ... |
| Gietz | 110 | Accounting |

19 rows selected.

Outer Joins Syntax

- You use an outer join to also see rows that do not meet the join condition.
- The outer join operator is the plus sign (+).

```
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column(+) = table2.column;
```

```
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column = table2.column(+);
```

Using Outer Joins to Return Records with No Direct Match

The missing rows can be returned if an outer join operator is used in the join condition. The operator is a plus sign enclosed in parentheses (+), and it is placed on the side of the join that is deficient in information. This operator has the effect of creating one or more null rows, to which one or more rows from the nondeficient table can be joined.

In the syntax:

| | |
|--------------------------------|---|
| <code>table1.column =</code> | is the condition that joins (or relates) the tables together. |
| <code>table2.column (+)</code> | is the outer join symbol, which can be placed on either side of the WHERE clause condition, but not on both sides. (Place the outer join symbol following the name of the column in the table without the matching rows.) |

Using Outer Joins

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE e.department_id(+) = d.department_id;
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Whalen | 10 | Administration |
| Hartstein | 20 | Marketing |
| Fay | 20 | Marketing |
| Mourgos | 50 | Shipping |
| Rajs | 50 | Shipping |
| Higgins | 110 | Accounting |
| Gietz | 110 | Accounting |
| | | Contracting |

20 rows selected

Outer Join Restrictions

- The outer join operator can appear on only *one* side of the expression: the side that has information missing. It returns those rows from one table that have no direct match in the other table.
- A condition involving an outer join cannot use the IN operator or be linked to another condition by the OR operator.

Self Joins

EMPLOYEES (WORKER)

| EMPLOYEE_ID | LAST_NAME | MANAGER_ID |
|-------------|-----------|------------|
| 101 | Kochhar | 100 |
| 102 | De Haan | 100 |
| 124 | Mourgos | 100 |
| 143 | Zlotkey | 100 |
| 201 | Hartstein | 100 |
| 200 | Whalen | 101 |
| 205 | Gietz | 201 |

19 rows selected

EMPLOYEES (MANAGER)

| EMPLOYEE_ID | LAST_NAME |
|-------------|-----------|
| 100 | King |
| 100 | King |
| 100 | King |
| 100 | King |
| 100 | King |
| 100 | King |
| 101 | Kochhar |
| 205 | Higgins |



**MANAGER_ID in the WORKER table is equal to
EMPLOYEE_ID in the MANAGER table.**

Sometimes you need to join a table to itself. To find the name of each employee's manager, you need to join the EMPLOYEES table to itself, or perform a self join. For example, to find the name of Whalen's manager, you need to:

- Find Whalen in the EMPLOYEES table by looking at the LAST_NAME column.
- Find the manager number for Whalen by looking at the MANAGER_ID column. Whalen's manager number is 101.
- Find the name of the manager with EMPLOYEE_ID 101 by looking at the LAST_NAME column. Kochhar's employee number is 101, so Kochhar is Whalen's manager.

In this process, you look in the table twice. The first time you look in the table to find Whalen in the LAST_NAME column and MANAGER_ID value of 101. The second time you look in the EMPLOYEE_ID column to find 101 and the LAST_NAME column to find Kochhar.

Joining a Table to Itself

```
SELECT worker.last_name || ' works for '
       || manager.last_name
FROM   employees worker, employees manager
WHERE  worker.manager_id = manager.employee_id;
```

| W.LAST_NAME 'WORKSFOR' M.LAST_NAME |
|--------------------------------------|
| Kochhar works for King |
| De Haan works for King |
| Mourgos works for King |
| Zlotkey works for King |
| Hartstein works for King |
| Whalen works for Kochhar |
| Higgins works for Kochhar |
| Usuald works for De Haan |
| Ray works for Hartstein |
| Gietz works for Higgins |

19 rows selected.

Joining Tables Using SQL: 1999 Syntax

Use a join to query data from more than one table.

```
SELECT  table1.column, table2.column
FROM    table1
[CROSS JOIN table2] |
[NATURAL JOIN table2] |
[JOIN table2 USING (column_name)] |
[JOIN table2
  ON(table1.column_name = table2.column_name)] |
[LEFT|RIGHT|FULL OUTER JOIN table2
  ON (table1.column_name = table2.column_name)];
```

Defining Joins

Using the SQL: 1999 syntax, you can obtain the same results as what was shown in the prior pages.

In the syntax:

| | |
|-----------------------------|--|
| <i>table1.column</i> | Denotes the table and column from which data is retrieved |
| CROSS JOIN | Returns a Cartesian product from the two tables |
| NATURAL JOIN | Joins two tables based on the same column name |
| JOIN <i>table</i> | |
| USING <i>column_name</i> | Performs an equijoin based on the column name |
| JOIN <i>table</i> ON | |
| <i>table1.column_name</i> | Performs an equijoin based on the condition in the ON clause |
| = <i>table2.column_name</i> | |
| LEFT/RIGHT/FULL OUTER | |

Creating Cross Joins

- The **CROSS JOIN** clause produces the cross-product of two tables.
- This is the same as a Cartesian product between the two tables.

```
SELECT last_name, department_name
FROM   employees
CROSS JOIN departments;
```

| LAST_NAME | DEPARTMENT_NAME |
|-----------|-----------------|
| Carstein | Accounting |
| Fay | Contracting |
| Higgins | Contracting |
| Gietz | Contracting |

150 rows selected.

Creating Natural Joins

- The **NATURAL JOIN** clause is based on all columns in the two tables that have the same name.
- It selects rows from the two tables that have equal values in all matched columns.
- If the columns having the same names have different data types, then an error is returned.

Note: The join can happen only on columns having the same names and data types in both the tables. If the columns have the same name, but different data types, then the **NATURAL JOIN** syntax causes an error.

Retrieving Records with Natural Joins

```
SELECT department_id, department_name,
       location_id, city
FROM   departments
NATURAL JOIN locations;
```

| DEPARTMENT_ID | DEPARTMENT_NAME | LOCATION_ID | CITY |
|---------------|-----------------|-------------|---------------------|
| 60 | IT | 1400 | Southlake |
| 50 | Shipping | 1500 | South San Francisco |
| 10 | Administration | 1700 | Seattle |
| 80 | Executive | 1700 | Seattle |
| 110 | Accounting | 1700 | Seattle |
| 190 | Contracting | 1700 | Seattle |
| 20 | Marketing | 1800 | Toronto |
| 90 | Sales | 2500 | Oxford |

Equijoins

The natural join can also be written as an equijoin:

```
SELECT department_id, department_name,  
       departments.location_id, city  
FROM   departments, locations  
WHERE  departments.location_id = locations.location_id;
```

Natural Joins with a **WHERE** Clause

Additional restrictions on a natural join are implemented by using a **WHERE** clause. The example below limits the rows of output to those with a department ID equal to 20 or 50.

```
SELECT  department_id, department_name,  
        location_id, city  
FROM    departments  
NATURAL JOIN locations  
WHERE   department_id IN (20, 50);
```

Creating Joins with the **USING** Clause

- **If several columns have the same names but the data types do not match, the **NATURAL JOIN** clause can be modified with the **USING** clause to specify the columns that should be used for an equijoin.**
Note: Use the **USING** clause to match only one column when more than one column matches.
- **Do not use a table name or alias in the referenced columns.**
- **The **NATURAL JOIN** and **USING** clauses are mutually exclusive.**

Natural joins use all columns with matching names and data types to join the tables. The **USING** clause can be used to specify only those columns that should be used for an equijoin. The columns referenced in the **USING** clause should not have a qualifier (table name or alias) anywhere in the SQL statement.

For example, this statement is valid:

```
SELECT l.city, d.department_name  
FROM   locations l JOIN departments d USING (location_id)  
WHERE  location_id = 1400;
```

This statement is invalid because the **LOCATION_ID** is qualified in the where clause:

```
SELECT l.city, d.department_name  
FROM locations l JOIN departments d USING (location_id)  
WHERE d.location_id = 1400;  
ORA-25154: column part of USING clause cannot have qualifier
```

The same restriction applies to **NATURAL** joins also. Therefore columns that have the same name in both tables have to be used without any qualifiers.

```
SELECT e.employee_id, e.last_name, d.location_id
FROM employees e JOIN departments d
USING (department_id);
```

| EMPLOYEE_ID | LAST_NAME | LOCATION_ID |
|-------------|-----------|-------------|
| 200 | Whalen | 1700 |
| 201 | Hartstein | 1800 |
| 202 | Fay | 1800 |
| 124 | Mourgos | 1500 |
| 141 | Rajs | 1500 |
| 142 | Davies | 1500 |
| 205 | Higgins | 1700 |
| 206 | Gietz | 1700 |

The example shown joins the DEPARTMENT_ID column in the EMPLOYEES and DEPARTMENTS tables, and thus shows the location where an employee works.

This can also be written as an equijoin:

```
SELECT employee_id, last_name,
       employees.department_id, location_id
FROM employees, departments
WHERE employees.department_id = departments.department_id;
```

Creating Joins with the ON Clause

- The join condition for the natural join is basically an equijoin of all columns with the same name.
- To specify arbitrary conditions or specify columns to join, the ON clause is used.
- Separates the join condition from other search conditions.
- The ON clause makes code easy to understand.

```
SELECT e.employee_id, e.last_name, e.department_id,
       d.department_id, d.location_id
FROM employees e JOIN departments d
ON (e.department_id = d.department_id);
```

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_ID | LOCATION_ID |
|-------------|-----------|---------------|---------------|-------------|
| 200 | Whalen | 10 | 10 | 1700 |
| 201 | Hartstein | 20 | 20 | 1800 |
| 202 | Fay | 20 | 20 | 1800 |
| 124 | Mourgos | 50 | 50 | 1500 |
| 141 | Rajs | 50 | 50 | 1500 |
| 142 | Davies | 50 | 50 | 1500 |
| 205 | Higgins | 110 | 110 | 1700 |
| 206 | Gietz | 110 | 110 | 1700 |

10 rows selected

The ON clause can also be used as follows to join columns that have different names:

```
SELECT e.last_name emp, m.last_name mgr
FROM   employees e JOIN employees m
ON     (e.manager_id = m.employee_id);
```

| EMP | MGR |
|---------|---------|
| Kochhar | King |
| De Haan | King |
| Mourgos | King |
| | |
| Gietz | Higgins |

19 rows selected.

The preceding example is a self join of the EMPLOYEE table to itself, based on the EMPLOYEE_ID and MANAGER_ID columns.

Creating Three-Way Joins with the ON Clause

```
SELECT employee_id, city, department_name
FROM   employees e
JOIN   departments d
ON     d.department_id = e.department_id
JOIN   locations l
ON     d.location_id = l.location_id;
```

| EMPLOYEE_ID | CITY | DEPARTMENT_NAME |
|-------------|---------------------|-----------------|
| 100 | Seattle | Executive |
| 101 | Seattle | Executive |
| 102 | Seattle | Executive |
| 103 | Southlake | IT |
| 104 | Southlake | IT |
| 107 | Southlake | IT |
| 124 | South San Francisco | Shipping |
| 206 | Seattle | IT |

19 rows selected

Three-Way Joins

A three-way join is a join of three tables. In SQL: 1999 compliant syntax, joins are performed from left to right, so the first join to be performed is EMPLOYEES JOIN DEPARTMENTS. The first join condition can reference columns in EMPLOYEES and DEPARTMENTS but cannot reference columns in LOCATIONS. The second join condition can reference columns from all three tables.

This can also be written as a three-way equijoin:

```
SELECT employee_id, city, department_name
FROM   employees, departments, locations
WHERE  employees.department_id = departments.department_id
AND    departments.location_id = locations.location_id;
```


INNER versus OUTER Joins

- In SQL: 1999, the join of two tables returning only matched rows is an inner join.
- A join between two tables that returns the results of the inner join as well as unmatched rows left (or right) tables is a left (or right) outer join.
- A join between two tables that returns the results of an inner join as well as the results of a left and right join is a full outer join.

Joins: Comparing SQL: 1999 to Oracle Syntax

| Oracle | SQL: 1999 |
|-------------------|-----------------------|
| Equijoin | Natural or Inner Join |
| Outerjoin | Left Outer Join |
| Selfjoin | Join ON |
| Nonequijoin | Join USING |
| Cartesian Product | Cross Join |

LEFT OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e
LEFT OUTER JOIN departments d
ON (e.department_id = d.department_id);
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| King | 90 | Executive |
| Kochhar | 90 | Executive |
| ... | ... | ... |
| Ernst | 60 | IT |
| Grant | | |
| Whalen | 10 | Administration |
| Hartstein | 20 | Marketing |
| Fay | 20 | Marketing |
| Higgins | 110 | Accounting |
| Gietz | 110 | Accounting |

20 rows selected.

This query retrieves all rows in the `EMPLOYEES` table, which is the left table even if there is no match in the `DEPARTMENTS` table.

This query was completed in earlier releases as follows:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE d.department_id (+) = e.department_id;
```

RIGHT OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e
RIGHT OUTER JOIN departments d
ON (e.department_id = d.department_id);
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Whalen | 10 | Administration |
| Hartstein | 20 | Marketing |
| Fay | 20 | Marketing |
| Mourgos | 50 | Shipping |
| Rajs | 50 | Shipping |
| Davies | 50 | Shipping |
| Matos | 50 | Shipping |
| | | |
| Gietz | 110 | Accounting |
| | | Contracting |

20 rows selected.

This query retrieves all rows in the DEPARTMENTS table, which is the right table even if there is no match in the EMPLOYEES table.

This query was completed in earlier releases as follows:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE d.department_id = e.department_id (+);
```

FULL OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e
FULL OUTER JOIN departments d
ON (e.department_id = d.department_id);
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Abel | 60 | Sales |
| Davies | 50 | Shipping |
| De Haan | 90 | Executive |
| Ernst | 60 | IT |
| Fay | 20 | Marketing |
| Gietz | 110 | Accounting |
| Grant | | |
| Hartstein | 20 | Marketing |
| | | |
| Zlotkey | 80 | Support |
| | | Contracting |

21 rows selected.

This query retrieves all rows in the EMPLOYEES table, even if there is no match in the DEPARTMENTS table. It also retrieves all rows in the DEPARTMENTS table, even if there is no match in the EMPLOYEES table.

Additional Conditions

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
FROM   employees e JOIN departments d  
ON     (e.department_id = d.department_id)  
AND    e.manager_id = 149;
```

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_ID | LOCATION_ID |
|-------------|-----------|---------------|---------------|-------------|
| 174 | Abel | 80 | 80 | 2500 |
| 176 | Taylor | 80 | 80 | 2500 |

You can apply additional conditions in the WHERE clause. The example shown performs a join on the EMPLOYEES and DEPARTMENTS tables, and, in addition, displays only employees with a manager ID equal to 149.

ASSIGNMENTS

- Write a query to display the last name, department number, and department name for all employees.

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Whalen | 10 | Administration |
| Hartstein | 20 | Marketing |
| Fay | 20 | Marketing |
| Mourgos | 50 | Shipping |
| Rajs | 50 | Shipping |
| Davies | 50 | Shipping |
| Matos | 50 | Shipping |
| Vargas | 50 | Shipping |
| De Haan | 50 | Shipping |
| Turner | 50 | Shipping |
| Baer | 50 | Shipping |
| Gietz | 110 | Accounting |

19 rows selected.

- Create a unique listing of all jobs that are in department 30. Include the location of department 90 in the output.

| JOB_ID | LOCATION_ID |
|--------|-------------|
| SA_MAN | 2500 |
| SA_REP | 2500 |

- Write a query to display the employee last name, department name, location ID, and city of all employees who earn a commission.

| LAST_NAME | DEPARTMENT_NAME | LOCATION_ID | CITY |
|-----------|-----------------|-------------|--------|
| Zlotkey | Sales | 2500 | Oxford |
| Abel | Sales | 2500 | Oxford |
| Taylor | Sales | 2500 | Oxford |

4. Display the employee last name and department name for all employees who have an *a* (lowercase) in their last names. Place your SQL statement in a text file named `lab4_4.sql`.

| LAST_NAME | DEPARTMENT_NAME |
|-----------|-----------------|
| Whalen | Administration |
| Hartstein | Marketing |
| Fay | Marketing |
| Rajs | Shipping |
| Davies | Shipping |
| Matos | Shipping |
| Vargas | Shipping |
| Taylor | Sales |
| Kochhar | Executive |
| De Haan | Executive |

10 rows selected.

5. Write a query to display the last name, job, department number, and department name for all employees who work in Toronto.

| LAST_NAME | JOB_ID | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|--------|---------------|-----------------|
| Hartstein | MK_MAN | 20 | Marketing |
| Fay | MK_REP | 20 | Marketing |

6. Display the employee last name and employee number along with their manager's last name and manager number. Label the columns *Employee*, *Emp#*, *Manager*, and *Mgr#*, respectively. Place your SQL statement in a text file named `lab4_6.sql`.

| Employee | EMP# | Manager | Mgr# |
|----------|------|---------|------|
| Kochhar | 101 | King | 100 |
| De Haan | 102 | King | 100 |
| Mourgos | 124 | King | 100 |
| Zlotkey | 149 | King | 100 |

| | | | |
|--------|-----|-----------|-----|
| Abel | 174 | Zlotkey | 149 |
| Taylor | 176 | Zlotkey | 149 |
| Grant | 178 | Zlotkey | 149 |
| Fay | 202 | Hartstein | 201 |
| Gietz | 206 | Higgins | 205 |

19 rows selected.

7. Modify lab4_6.sql to display all employees including King, who has no manager. Order the results by the employee number.

Place your SQL statement in a text file named lab4_7.sql. Run the query in lab4_7.sql.

| Employee | EMP# | Manager | Mgr# |
|----------|------|---------|------|
| King | 100 | | |
| Kochhar | 101 | King | 100 |
| De Haan | 102 | King | 100 |
| Hunold | 103 | De Haan | 102 |
| Ernst | 104 | Hunold | 103 |
| Lorentz | 107 | Hunold | 103 |
| Mourgos | 124 | King | 100 |
| Higgins | 205 | Kochhar | 101 |
| Gietz | 206 | Higgins | 205 |

20 rows selected.

8. Create a query that displays employee last names, department numbers, and all the employees who work in the same department as a given employee. Give each column an appropriate label.

| DEPARTMENT | EMPLOYEE | COLLEAGUE |
|------------|-----------|-----------|
| 20 | Fay | Hartstein |
| 20 | Hartstein | Fay |
| 50 | Davies | Matos |
| 50 | Davies | Mourgos |
| 50 | Davies | Rajs |
| 50 | Davies | Vargas |
| 50 | Matos | Davies |
| 50 | Matos | Mourgos |
| 50 | Matos | Rajs |
| 50 | Matos | Vargas |
| 110 | Gietz | Higgins |
| 110 | Higgins | Gietz |

42 rows selected.

9. Show the structure of the `JOB_GRADES` table. Create a query that displays the name, job, department name, salary, and grade for all employees.

| Name | Null? | Type |
|-------------|-------|-------------|
| GRADE_LEVEL | | VARCHAR2(3) |
| LOWEST_SAL | | NUMBER |
| HIGHEST_SAL | | NUMBER |

| LAST_NAME | JOB_ID | DEPARTMENT_NAME | SALARY | GRA |
|-----------|----------|-----------------|--------|-----|
| Matos | ST_CLERK | Shipping | 2600 | A |
| Vargas | ST_CLERK | Shipping | 2500 | A |
| Lorentz | IT_PROG | IT | 4200 | B |
| Mourgos | ST_MAN | Shipping | 5800 | B |
| Rajs | ST_CLERK | Shipping | 3500 | B |
| Davies | ST_CLERK | Shipping | 3100 | B |
| Whalen | AD_ASST | Administration | 4400 | B |

| | | | | |
|---------|-------|-----------|-------|---|
| De Haan | AD_VP | Executive | 17000 | E |
|---------|-------|-----------|-------|---|

19 rows selected.

10. Create a query to display the name and hire date of any employee hired after employee Davies.

| LAST_NAME | HIRE_DATE |
|-----------|-----------|
| Lorentz | 07-FEB-99 |
| Mourgos | 16-NOV-99 |
| Matos | 15-MAR-98 |
| Vargas | 09-JUL-98 |
| Zlotkey | 29-JAN-00 |
| Taylor | 24-MAR-98 |
| Grant | 24-MAY-99 |
| Fay | 17-AUG-97 |

8 rows selected.

11. Display the names and hire dates for all employees who were hired before their managers, along with their manager's names and hire dates. Label the columns Employee, Emp Hired, Manager, and Mgr Hired, respectively.

| Employee | Emp Hired | Manager | Mgr Hired |
|----------|-----------|---------|-----------|
| Whalen | 17-SEP-87 | Kochhar | 21-SEP-89 |
| Hunold | 03-JAN-90 | De Haan | 13-JAN-93 |
| Rajs | 17-OCT-95 | Mourgos | 16-NOV-99 |
| Davies | 29-JAN-97 | Mourgos | 16-NOV-99 |
| Matos | 15-MAR-98 | Mourgos | 16-NOV-99 |
| Vargas | 09-JUL-98 | Mourgos | 16-NOV-99 |
| Abel | 11-MAY-96 | Zlotkey | 29-JAN-00 |
| Taylor | 24-MAR-98 | Zlotkey | 29-JAN-00 |
| Grant | 24-MAY-99 | Zlotkey | 29-JAN-00 |

9 rows selected.